


SQL



DATABASES, VIEWS, AND TABLES

DATABASES, VIEWS, AND TABLES

Besides the DataFrame API, spark implements the concepts of databases, functions, views and tables. The idea is to enable us to use standard SQL to process data

Spark stores metadata concerning those elements in an internal catalog, which is accessible from a SparkSession

```
[100]: spark.catalog.listTables()
[100]: [Table(nam f listColumns function
[21]: spark.sql(f listDatabases function
[21]: DataFrame[ f listFunctions function
[21]: DataFrame[ f listTables function
[23]: usersDf.wr f recoverPartitions function
[25]: spark.cata f refreshByPath function
[25]: [Table(nam f refreshTable function
[25]: [Table(nam f registerFunction function
[25]: [Table(nam i setCurrentDatabase instance
[25]: [Table(nam f uncacheTable function
Table(name= pessoas , database=None, descriptio
```

CREATING A VIEW

```
# Para carregar um DataFrame diretamente de uma fonte de dados externa, pode-se usar:  
df = spark.sql("SELECT * FROM csv.`users.csv`",)  
df.printSchema()  
df.show()
```

```
root  
 |-- _c0: string (nullable = true)  
 |-- _c1: string (nullable = true)  
 |-- _c2: string (nullable = true)
```

```
+---+-----+  
|_c0|          _c1|_c2|  
+---+-----+  
| 1| 'Fabio Nogueira'| 47|  
| 2| 'Andrea Vasconcelos'| 47|  
| 3| 'Thiago Vasconcel...| 21|  
+---+-----+
```

```
# Não há ainda tabelas, uma vez que trouxemos direto o conteúdo do arquivo para um dataframe  
catalog.listTables()
```

```
[]
```

```
# Criando uma tabela temporária a partir de um dataframe  
df.createOrReplaceTempView('pessoas')
```

```
catalog.listTables()
```

```
[Table(name='pessoas', database=None, description=None, tableType='TEMPORARY', isTemporary=True)]
```

QUERYING VIEWS USING SQL

```
# Creating a temporary table to enable us to submit sql directly  
users.createOrReplaceTempView('pessoas')
```

```
pes = spark.sql('select * from pessoas')  
pes.show()
```

```
+---+-----+  
|_c0|          _c1|_c2|  
+---+-----+  
|  1|    'Fabio Nogueira'| 47|  
|  2| 'Andrea Vasconcelos'| 47|  
|  3| 'Thiago Vasconcel...'| 18|  
+---+-----+
```

TEMPORARY VIEWS

Temporary views in Spark SQL are session-scoped and will disappear if the session that creates it terminates

If you want to have a temporary view that is shared among all sessions and keep alive until the Spark application terminates, you can create a global temporary view

TEMPORARY VIEWS

```
# Register the DataFrame as a global temporary view
df.createGlobalTempView("people")

# Global temporary view is tied to a system preserved database 'global_temp'
spark.sql("SELECT * FROM global_temp.people").show()
# +-----+-----+
# | age|   name|
# +-----+-----+
# |null|Michael|
# | 30|   Andy|
# | 19|  Justin|
# +-----+-----+

# Global temporary view is cross-session
spark.newSession().sql("SELECT * FROM global_temp.people").show()
# +-----+-----+
# | age|   name|
# +-----+-----+
# |null|Michael|
# | 30|   Andy|
# | 19|  Justin|
# +-----+-----+
```

PERSISTING DATAFRAMES

DataFrames can also be saved as persistent tables into Hive metastore using the `saveAsTable` command

These tables can be organized in Databases

Unlike the `createOrReplaceTempView` command, `saveAsTable` will materialize the contents of the DataFrame and create a pointer to the data in the Hive metastore.

A DataFrame for a persistent table can be created by calling the `table` method on a `SparkSession` with the name of the table.

PERSISTING DATAFRAMES

For file-based data source, e.g. text, parquet, json, etc. you can specify a custom table path via the path option, e.g. `df.write.option("path", "/some/path").saveAsTable("t")`. Table persisted that way are referred to as “External”, this means that when the table is dropped, the custom table path will not be removed and the table data is still there. If no custom table path is specified, we say the table is “Managed” by Spark. So, it will write data to a default table path under the warehouse directory. When the table is dropped, the default table path will be removed too.

Persistent datasource tables have per-partition metadata stored in the Hive metastore.

```
# 1/24/M/technician/85711
usersDf = spark.read.csv(path='/home/jovyan/work/ml-100k/u.user',sep='|',schema='id int, age int, genre string, occupation string, time int')
usersDf.show(5)
usersDf.printSchema()
```

```
+---+---+-----+-----+
| id|age|genre|occupation| time|
+---+---+-----+-----+
|  1| 24|   M|technician|85711|
|  2| 53|   F|   other|94043|
|  3| 23|   M|   writer|32067|
|  4| 24|   M|technician|43537|
|  5| 33|   F|   other|15213|
+---+---+-----+-----+
only showing top 5 rows
```

```
root
 |-- id: integer (nullable = true)
 |-- age: integer (nullable = true)
 |-- genre: string (nullable = true)
 |-- occupation: string (nullable = true)
 |-- time: integer (nullable = true)
```

```
spark.sql("create database ml100k")
```

```
DataFrame[]
```

```
spark.catalog.listTables()
```

```
[Table(name='pessoas', database=None, description=None, tableType='TEMPORARY', isTemporary=True)]
```

```
spark.sql('DROP TABLE IF EXISTS ml100k.users')
```

```
DataFrame[]
```


```
usersDf.write.saveAsTable('ml100k.users')
```

```
spark.catalog.listTables('ml100k')
```

```
[Table(name='users', database='ml100k', description=None, tableType='MANAGED', isTemporary=False),
 Table(name='pessoas', database=None, description=None, tableType='TEMPORARY', isTemporary=True)]
```



/ ... / spark-warehouse / ml100k.db /	
Name	Last Modified
ages	2 hours ago
users	5 hours ago



FUNCTIONS




FUNCTIONS

Structured APIs includes lots of built-in functions that can be applied to columns according to their types

We will show some of them, others can be found here:

- <https://spark.apache.org/docs/latest/api/python/pyspark.sql.html#module-pyspark.sql.functions>

Moreover, Spark allows users to define their own functions, which can be registered and used in their queries



NUMERICAL



GENERATING IDS

```
: ratingsDf.printSchema()
```

```
root
|-- userid: integer (nullable = true)
|-- itemid: integer (nullable = true)
|-- rating: integer (nullable = true)
|-- time: integer (nullable = true)
```

```
: ratingsWithKey = ratingsDf.select(monotonically_increasing_id().alias('id'),'*')
ratingsWithKey.printSchema()
```

```
root
|-- id: long (nullable = false)
|-- userid: integer (nullable = true)
|-- itemid: integer (nullable = true)
|-- rating: integer (nullable = true)
|-- time: integer (nullable = true)
```

```
: ratingsWithKey.show(5)
```

```
+---+-----+-----+-----+
| id|userid|itemid|rating|   time|
+---+-----+-----+-----+
|  0|  196|  242|    3|881250949|
|  1|  186|  302|    3|891717742|
|  2|   22|  377|    1|878887116|
|  3|  244|   51|    2|880606923|
|  4|  166|  346|    1|886397596|
+---+-----+-----+-----+
only showing top 5 rows
```


POW

```
1 from pyspark.sql.functions import expr, pow
2 fabricatedQuantity = pow(col("Quantity") * col("UnitPrice"), 2) + 5
3 df.select(expr("CustomerId"), fabricatedQuantity.alias("realQuantity")).show(2)
```

```
1 df.selectExpr(
2     "CustomerId",
3     "(POWER((Quantity * UnitPrice), 2.0) + 5) as realQuantity").show(2)
```

ROUND AND BROUND

```
1 from pyspark.sql.functions import *  
2 df.select(round(col("UnitPrice"), 1).alias("rounded"), col("UnitPrice")).show(5)
```



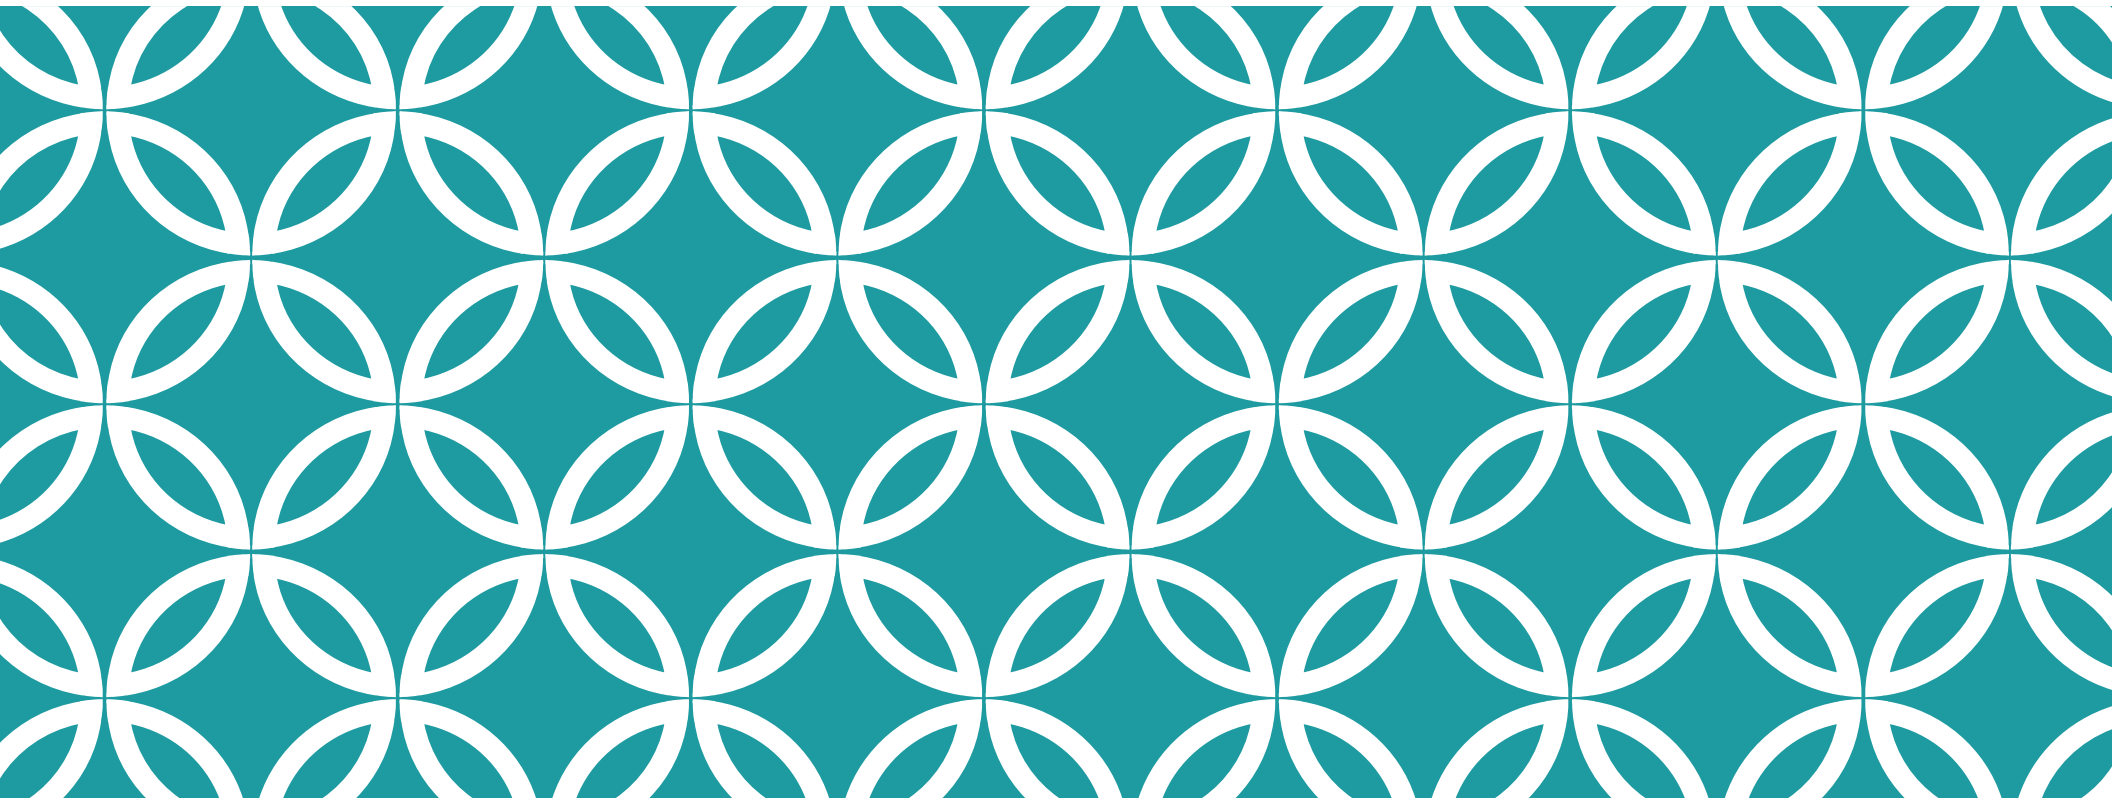
STATISTICS

```
1 df.describe().show()
2 df.describe('UnitPrice').show()
```



```
count, mean, stddev_pop, min, max
```





STRINGS |

CASE

To modify a string column in order to change the case of its contents, we can use *initcap*, *lower*, and *upper* functions

```
1 from pyspark.sql.functions import col, initcap, lower, upper
2 df.select(col("Description"), \
3         initcap(col("Description")), \
4         lower(col("Description")), \
5         upper(col("Description"))).show()
```

```
1 +-----+-----+-----+-----+
2 |      Description|initcap(Description)|  lower(Description)|  upper(Description)|
3 +-----+-----+-----+-----+
4 |WHITE HANGING HEA...|White Hanging Hea...|white hanging hea...|WHITE HANGING HEA..
5 | WHITE METAL LANTERN| White Metal Lantern| white metal lantern| WHITE METAL LANTER
6 |CREAM CUPID HEART...|Cream Cupid Heart...|cream cupid heart...|CREAM CUPID HEART..
7 |KNITTED UNION FLA...|Knitted Union Fla...|knitted union fla...|KNITTED UNION FLA..
8 |RED WOOLLY HOTTIE...|Red Woolly Hottie...|red woolly hottie...|RED WOOLLY HOTTIE..
```

DEALING WITH SPACES

```
from pyspark.sql.functions import lit, ltrim, rtrim, rpad, lpad, trim
df.select(
    ltrim(lit("  HELLO  ")).alias("ltrim"),
    rtrim(lit("  HELLO  ")).alias("rtrim"),
    trim(lit("  HELLO  ")).alias("trim"),
    lpad(lit("HELLO"), 3, "#").alias("lp"),
    rpad(lit("HELLO"), 10, "#").alias("rp")).show()
```

ltrim	rtrim	trim	lp	rp
HELLO	HELLO	HELLO	HEL	HELLO#####
HELLO	HELLO	HELLO	HEL	HELLO#####
HELLO	HELLO	HELLO	HEL	HELLO#####

But, wait!!! Why do I see 3 lines here???

```
df.show()
```

_c0	_c1	_c2
1	'Fabio Nogueira'	47
2	'Andrea Vasconcelos'	47
3	'Thiago Vasconcel...	25

REPLACING STRINGS

```
1 from pyspark.sql.functions import regexp_replace, col
2
3 regex1 = 'BLACK|WHITE|RED|GREEN|BLUE'
4 regex2 = r'\bBLACK\b|\bWHITE\b|\bRED\b|\bGREEN\b|\bBLUE\b'
5 regex3 = r'((?i)\bBLACK\b|\bWHITE\b|\bRED\b|\bGREEN\b|\bBLUE\b)'
6
7 mList = [('BLUEBLUE',), ('REDO it',), ('BLACK Black line',)]
8
9 memoryDf = spark.createDataFrame(mList, ['Desc',])
10
11 replacedDf = memoryDf.select('Desc',
12                               regexp_replace(col('Desc'), regex1, 'COLOR').alias('Patern 1'),
13                               regexp_replace(col('Desc'), regex2, 'COLOR').alias('Pattern 2'),
14                               regexp_replace(col('Desc'), regex3, 'COLOR').alias('Pattern 3'))
15
16 replacedDf.show(truncate=False)
```

```
1 +-----+-----+-----+-----+
2 |Desc          |Patern 1      |Pattern 2     |Pattern 3     |
3 +-----+-----+-----+-----+
4 |BLUEBLUE      |COLORCOLOR    |BLUEBLUE      |BLUEBLUE      |
5 |REDO it       |COLORO it     |REDO it       |REDO it       |
6 |BLACK Black line|COLOR Black line|COLOR Black line|COLOR COLOR line|
7 +-----+-----+-----+-----+
8
```

MODIFYING INDIVIDUAL CHARACTERS


```
1 from pyspark.sql.functions import translate
2 df.select(translate(col("Description"), "LEET", "1337"), col("Description"))\
3     .show(2)
```

```
1 +-----+-----+
2 |translate(Description, LEET, 1337)|      Description|
3 +-----+-----+
4 |          WHI73 HANGING H3A...|WHITE HANGING HEA...|
5 |          WHI73 M37A1 1AN73RN| WHITE METAL LANTERN|
6 +-----+-----+
```

VERIFYING IF A STRING IS IN A COLUMN

```
1 from pyspark.sql.functions import instr
2 containsBlack = instr(col("Description"), "BLACK") >= 1
3 containsWhite = instr(col("Description"), "WHITE") >= 1
4 df.withColumn("hasSimpleColor", containsBlack | containsWhite)\
5   .where("hasSimpleColor")\
6   .select("Description").show(3, False)
```

```
1 +-----+
2 |Description|
3 +-----+
4 |WHITE HANGING HEART T-LIGHT HOLDER|
5 |WHITE METAL LANTERN|
6 |RED WOOLLY HOTTIE WHITE HEART.|
7 +-----+
```



DATES AND TIME



CURRENT DATE AND TIMESTAMP

```
1 from pyspark.sql.functions import current_date, current_timestamp
2 # spark.range(start, end=None, step=1, numPartitions=None)
3 # Create a DataFrame with single pyspark.sql.types.LongType column named id, contain
4
5 dateDF = spark.range(10)\
6     .withColumn("today", current_date())\
7     .withColumn("now", current_timestamp())
8 dateDF.createOrReplaceTempView("dateTable")
9 dateDF.show(truncate=False)
10 dateDF.printSchema()
```

```
1 +---+-----+-----+
2 |id|today      |now                |
3 +---+-----+-----+
4 |0 |2019-10-31|2019-10-31 14:54:29.498|
5 |1 |2019-10-31|2019-10-31 14:54:29.498|
6 |2 |2019-10-31|2019-10-31 14:54:29.498|
7 |3 |2019-10-31|2019-10-31 14:54:29.498|
8 |4 |2019-10-31|2019-10-31 14:54:29.498|
9 |5 |2019-10-31|2019-10-31 14:54:29.498|
10 |6 |2019-10-31|2019-10-31 14:54:29.498|
11 |7 |2019-10-31|2019-10-31 14:54:29.498|
12 |8 |2019-10-31|2019-10-31 14:54:29.498|
13 |9 |2019-10-31|2019-10-31 14:54:29.498|
14 +---+-----+-----+
```

ADD AND SUB DATES

```
1 from pyspark.sql.functions import col, date_add, date_sub
2 dateDF.select('id', 'today', date_sub(col("today"), 5), date_add(col("today"), 5)).
```

```
1 +---+-----+-----+-----+
2 | id|      today|date_sub(today, 5)|date_add(today, 5)|
3 +---+-----+-----+-----+
4 |  0|2019-10-31|      2019-10-26|      2019-11-05|
5 +---+-----+-----+-----+
```

```

1 from pyspark.sql.functions import datediff, months_between, to_date, lit
2 dateDF.withColumn("week_ago", date_sub(col("today"), 7)) \
3   .withColumn('next_week', date_add(col('today'),7)) \
4   .select(datediff(col("week_ago"), col("today")).alias('DaysFromWeekAgo'), \
5           datediff(col('next_week'), col('today')).alias('DaysFromNextWeek')).show(
6
7 dateDF.select(
8     to_date(lit("2016-01-01")).alias("start"),
9     to_date(lit("2017-05-22")).alias("end"))\
10  .select(months_between(col("start"), col("end"))).show(1)

```

```

1 +-----+-----+
2 |DaysFromWeekAgo|DaysFromNextWeek|
3 +-----+-----+
4 |           -7|           7|
5 +-----+-----+
6 only showing top 1 row
7
8 +-----+
9 |months_between(start, end, true)|
10 +-----+
11 |           -16.67741935|
12 +-----+
13 only showing top 1 row


```

CREATING DATES

We can convert a string or timestamp column to a *date* using *to_date* function. Besides the column, it can receive a format, specified according to *SimpleDateFormat* (<https://docs.oracle.com/javase/tutorial/i18n/format/simpleDateFormat.html>). If a value in the string column can not be recognized as a *date*, Spark will just return *null* as the corresponding value.

```
1 from pyspark.sql.functions import to_date
2 dateFormat = "yyyy-dd-MM"
3 cleanDateDF = spark.range(1).select(
4     to_date(lit("2017-12-11"), dateFormat).alias("date"),
5     to_date(lit("2017-20-12"), dateFormat).alias("date2")).show()
```

```
1 from pyspark.sql.functions import to_timestamp
2 cleanDateDF.select(to_timestamp(col("date"), dateFormat)).show()
```



WORKING WITH NULL



COALESCE

```
1 cDf = spark.createDataFrame([(None, None), (1, None), (None, 2)], ("a", "b"))
2 cDf.select('*',
3     lit('0.0').alias('lit(0.0)'),
4     coalesce(cDf["a"], cDf["b"]),
5     coalesce(cDf["a"], lit(0.0)),
6     coalesce(lit('0.0'), cDf['a'])
7 ).show()
```

	a	b	lit(0.0)	coalesce(a, b)	coalesce(a, 0.0)	coalesce(0.0, a)
	null	null	0.0	null	0.0	0.0
	1	null	0.0	1	1.0	0.0
	null	2	0.0	2	0.0	0.0

DROP

```
1 cDf = spark.createDataFrame([(None, None), (1, None), (None, 2), (1,1)], ("a", "b"))
2 cDf.show()
3 +-----+-----+
4 |    a|    b|
5 +-----+-----+
6 |null|null|
7 |    1|null|
8 |null|    2|
9 |    1|    1|
10 +-----+-----+
11 cDf.dropna().show()
12 +-----+
13 |    a|    b|
14 +-----+
15 |    1|    1|
16 +-----+
17 cDf.na.drop().show()
18 +-----+
19 |    a|    b|
20 +-----+
21 |    1|    1|
22 +-----+
23 cDf.na.drop('any').show()
24 +-----+
25 |    a|    b|
26 +-----+
27 |    1|    1|
28 +-----+
```

DROP (CONT)

```
29
30 cDf.na.drop('all').show()
31 +----+----+
32 |  a|  b|
33 +----+----+
34 |  1|null|
35 |null|  2|
36 |  1|  1|
37 +----+----+
38 cDf.na.drop(subset=['a']).show()
39 +----+----+
40 |  a|  b|
41 +----+----+
42 |  1|null|
43 |  1|  1|
44 +----+----+
45 cDf.na.drop(subset=['a', 'b']).show()
46 +----+----+
47 |  a|  b|
48 +----+----+
49 |  1|  1|
50 +----+----+
```


FILL

```
1 df4.show()
2 +----+-----+-----+
3 | age|height| name|
4 +----+-----+-----+
5 | 10|    80|Alice|
6 |  5|   null|  Bob|
7 |null|  null|  Tom|
8 |null|  null| null|
9 +----+-----+-----+
```



```
1 df4.na.fill(50).show()
2 +---+-----+-----+
3 |age|height| name|
4 +---+-----+-----+
5 | 10|    80|Alice|
6 |  5|    50|  Bob|
7 | 50|    50|  Tom|
8 | 50|    50| null|
9 +---+-----+-----+
```



FILL (CONT)

```
1 df4.na.fill(50, subset=['age']).show()
2 +---+-----+-----+
3 |age|height| name|
4 +---+-----+-----+
5 | 10|    80|Alice|
6 |  5|   null| Bob|
7 | 50|   null| Tom|
8 | 50|   null| null|
9 +---+-----+-----+
```

```
1 df4.na.fill({'age': 50, 'name': 'unknown'}).show()
2 +---+-----+-----+
3 |age|height|  name|
4 +---+-----+-----+
5 | 10|    80| Alice|
6 |  5|   null|  Bob|
7 | 50|   null|  Tom|
8 | 50|   null|unknown|
9 +---+-----+-----+
```

exit: Ctrl+↵

REPLACE

```
1 df4.show()
2 +----+-----+-----+
3 | age|height| name|
4 +----+-----+-----+
5 | 10|    80|Alice|
6 |  5|   null| Bob|
7 |null| null| Tom|
8 |null| null| null|
9 +----+-----+-----+
```




```
1 df4.na.replace(10, 20).show()
2 +----+-----+-----+
3 | age|height| name|
4 +----+-----+-----+
5 | 20|    80|Alice|
6 |  5|   null| Bob|
7 |null| null| Tom|
8 |null| null| null|
9 +----+-----+-----+
```



REPLACE (CONT)

```
1 df4.na.replace('Alice', None).show()
2 +----+-----+----+
3 | age|height|name|
4 +----+-----+----+
5 | 10|    80|null|
6 |  5|   null| Bob|
7 |null|  null| Tom|
8 |null|  null|null|
9 +----+-----+----+
```

```
1 df4.na.replace({10 : 100, 5 : 50, 80 : 10},subset=['age']).show()
2 +----+-----+----+
3 | age|height| name|
4 +----+-----+----+
5 | 100|    80|Alice|
6 |  50|   null| Bob|
7 |null|  null| Tom|
8 |null|  null| null|
9 +----+-----+----+
```



COMPLEX TYPES



STRUCTS

```
1 from pyspark.sql.functions import struct
2 complexDF = df.select(struct("Description", "InvoiceNo").alias("complex"))
3 complexDF.createOrReplaceTempView("complexDF")
4
5 complexDF.select("complex.Description")
6 complexDF.select(col("complex").getField("Description"))
7 complexDF.select("complex.*")
```



ARRAYS


```
1 from pyspark.sql.functions import split, size, array_contains
2
3 df.select('Description',
4           split(col("Description"), " ").alias("asArray"),
5           array_contains(split(col("Description"), " "), "WHITE"),
6           size(split(col("Description"), " ").alias('wordsCounter'))
7 ).show(2, truncate=False)
```

Description	asArray	array_contains
WHITE HANGING HEART T-LIGHT HOLDER	[WHITE, HANGING, HEART, T-LIGHT, HOLDER]	true
WHITE METAL LANTERN	[WHITE, METAL, LANTERN]	true

ARRAYS

 We can explode an array column to generate a line per element:

```
1 from pyspark.sql.functions import split, explode
2
3 df.withColumn("splitted", split(col("Description"), " "))\
4   .withColumn("exploded", explode(col("splitted"))) \
5   .withColumn('size',size("splitted")) \
6   .select("Description", "InvoiceNo", "exploded").show(5)
7 +-----+-----+-----+
8 |      Description|InvoiceNo|exploded|
9 +-----+-----+-----+
10 |WHITE HANGING HEA...|  536365|  WHITE|
11 |WHITE HANGING HEA...|  536365| HANGING|
12 |WHITE HANGING HEA...|  536365|  HEART|
13 |WHITE HANGING HEA...|  536365| T-LIGHT|
14 |WHITE HANGING HEA...|  536365|  HOLDER|
15 +-----+-----+-----+
```

AGGREGATION



COUNT

```
1 from pyspark.sql.functions import count
2
3 df4.show()
4 +-----+-----+-----+
5 | age|height| name|
6 +-----+-----+-----+
7 | 10|    80|Alice|
8 |  5|   null|  Bob|
9 |null|  null|  Tom|
10 |null|  null| null|
11 +-----+-----+-----+
12 df4.select(count('*'),
13            count("age"),
14            count("height"),
15            count('name')).show()
16 +-----+-----+-----+-----+
17 | count(1)|count(age)|count(height)|count(name)|
18 +-----+-----+-----+-----+
19 |         4|         2|         1|         3|
20 +-----+-----+-----+-----+
```

COUNT

```
1 from pyspark.sql.functions import countDistinct
2
3 dfx = spark.sparkContext.parallelize([(2, 'Alice'), (2, 'Bob')])\
4     .toDF(StructType([StructField('age', IntegerType()),
5                          StructField('name', StringType())]))
6
7 dfx.show()
8 +----+-----+
9 |age| name|
10 +----+-----+
11 | 2|Alice|
12 | 2|  Bob|
13 +----+-----+
14 dfx.agg(countDistinct(dfx.age).alias('c')).show()
15 +----+
16 |  c|
17 +----+
18 |  1|
19 +----+
20 dfx.agg(countDistinct(dfx.name).alias('c')).show()
21 +----+
22 |  c|
23 +----+
24 |  2|
25 +----+
26 dfx.agg(countDistinct(dfx.age, dfx.name).alias('c')).show()
27 +----+
28 |  c|
29 +----+
30 |  2|
31 +----+
```

FIRST AND LAST

```
: df.show(truncate=False)
```

```
+---+-----+---+
|uid|name                |age|
+---+-----+---+
|1  |'Fabio Nogueira'      |47 |
|2  |'Andrea Vasconcelos'  |47 |
|3  |'Thiago Vasconcelos Nogueira'|25 |
+---+-----+---+
```

```
: from pyspark.sql.functions import first, last
df.select(first('uid'), last('age')).show()
```

```
+-----+-----+
|first(uid, false)|last(age, false)|
+-----+-----+
|1|25|
+-----+-----+
```

SUM, MIN, AND MAX

```
1 from pyspark.sql.functions import sum, min, max
2 dfx.show()
3 +---+-----+
4 | age| name|
5 +---+-----+
6 |  2|Alice|
7 |  5|  Bob|
8 |null| Fool|
9 |null| null|
10 +---+-----+
11 dfx.select(sum("age"),
12           min("age"),
13           max("age")).show()
14 +-----+-----+-----+
15 |sum(age)|min(age)|max(age)|
16 +-----+-----+-----+
17 |      7|      2|      5|
18 +-----+-----+-----+
```

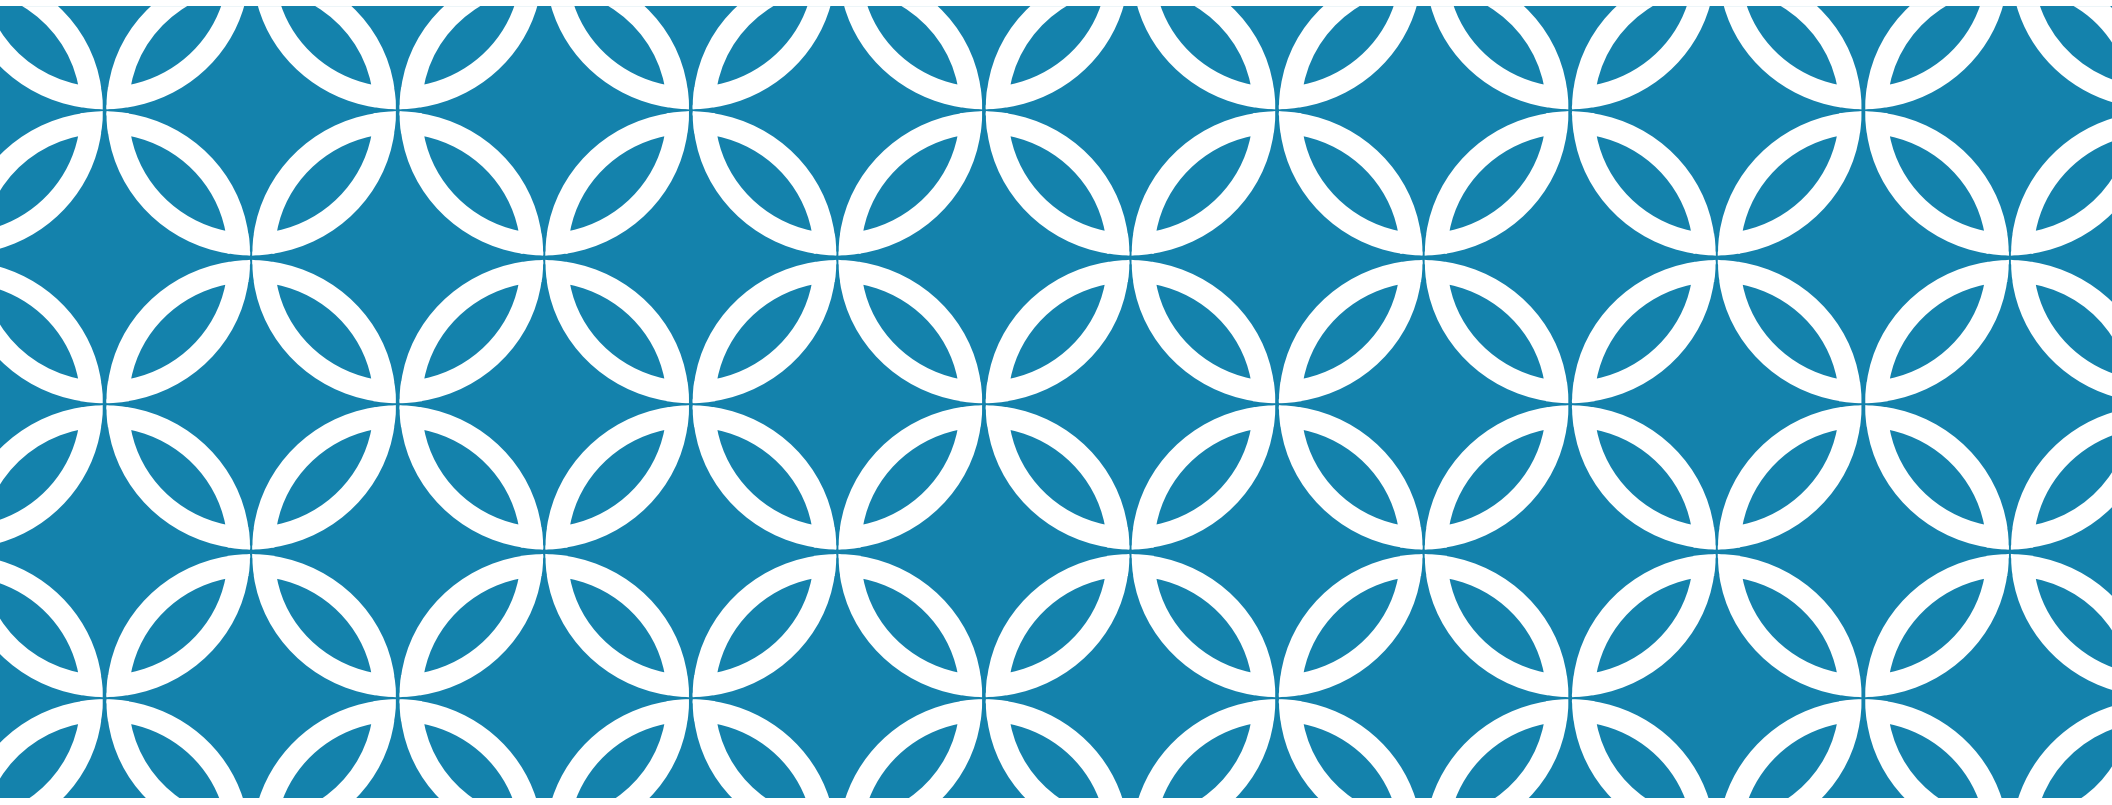
AVG, VAR, AND STDEV

```
1 dfx.select(  
2     count("age"),  
3     sum("age"),  
4     sum('age') / count('age') ,  
5     avg("age")).show()  
6 +-----+-----+-----+-----+  
7 |count(age)|sum(age)|(sum(age) / count(age))|avg(age)|  
8 +-----+-----+-----+-----+  
9 |         2|        7|              3.5|      3.5|  
10 +-----+-----+-----+-----+
```

```
1 from pyspark.sql.functions import var_pop, stddev_pop  
2 from pyspark.sql.functions import var_samp, stddev_samp  
3  
4 df.select(var_pop("Quantity"), var_samp("Quantity"),  
5     stddev_pop("Quantity"), stddev_samp("Quantity")).show()  
6 +-----+-----+-----+-----+  
7 |var_pop(Quantity)|var_samp(Quantity)|stddev_pop(Quantity)|stddev_samp(Quantity)|  
8 +-----+-----+-----+-----+  
9 |695.2492099104054| 695.4729785650273| 26.367578764657278| 26.371821677029203|  
10 +-----+-----+-----+-----+
```

GROUP BY

```
1 from pyspark.sql.functions import *
2 usersDf.groupBy('genre').count().\
3   na.replace({'F':'Female','M':'Male'}).show()
4 +-----+-----+
5 | genre|count|
6 +-----+-----+
7 |Female|  273|
8 |  Male|  670|
9 +-----+-----+
```



QUESTIONS???

